

第3课 vLLM部署

官方文档：https://docs.vllm.ai/en/latest/getting_started/quickstart/

▼ Installation Plain Text |

```
1 uv venv --python 3.12 --seed
2 source .venv/bin/activate
3 uv pip install vllm --torch-backend=auto
```

维度	Offline Batched Inference	Online Serving
使用入口	Python LLM 或 vllm run-batch	vllm serve
请求来源	本地程序、数据集、JSONL	HTTP 客户端、业务服务
请求到达	通常一次提供一批	随时动态到达
调用方式	同步、阻塞式为主	异步、并发式
结果返回	通常全部完成后返回列表	每个请求独立返回，可流式输出
Batch 组成	从已有请求集合中调度	随着请求到达和结束持续变化
服务生命周期	任务结束后进程通常退出	长期运行、模型常驻显存
主要目标	数据集整体吞吐量	吞吐量、首 Token 延迟和服务稳定性
工程能力	相对简单	HTTP、并发、取消、流控、监控、鉴权

Offline Batched Inference

▼ Offline Batched Inference

Python |

```
1  from vllm import LLM, SamplingParams
2
3  # Sample prompts.
4  prompts = [
5      "Hello, my name is",
6      "The president of the United States is",
7      "The capital of France is",
8      "The future of AI is",
9  ]
10 # Create a sampling params object.
11 sampling_params = SamplingParams(temperature=0.8, top_p=0.95)
12
13
14 def main():
15     # Create an LLM.
16     llm = LLM(model="facebook/opt-125m")
17     # Generate texts from the prompts.
18     # The output is a list of RequestOutput objects
19     # that contain the prompt, generated text, and other information.
20     outputs = llm.generate(prompts, sampling_params)
21     # Print the outputs.
22     print("\nGenerated Outputs:\n" + "-" * 60)
23     for output in outputs:
24         prompt = output.prompt
25         generated_text = output.outputs[0].text
26         print(f"Prompt:      {prompt!r}")
27         print(f"Output:       {generated_text!r}")
28         print("-" * 60)
29
30
31 if __name__ == "__main__":
32     main()
```

Generated Outputs:

```
-----
Prompt:   'Hello, my name is'
Output:   " Michael. I'm a manager of The Prudential Financial Trust, based in"
-----
Prompt:   'The president of the United States is'
Output:   ' allegedly planning to nominate a right-wing populist to replace the president of the United'
-----
Prompt:   'The capital of France is'
Output:   ' starting to pull back from a new order to cut the number of flights it says'
-----
Prompt:   'The future of AI is'
Output:   ' a terrifying picture\nEverything is on the line in the future of AI.\n'
-----
```

```
1 VLLM_USE_MODELSCOPE=True \  
2 HF_HUB_OFFLINE=1 \  
3 /root/.venv/bin/python -u /root/offline.py  
4  
5 VLLM_USE_MODELSCOPE=True 告诉 vLLM 优先通过 ModelScope 查找或下载模型，而不是 Hu  
gging Face。  
6 HF_HUB_OFFLINE=1 只使用本地缓存，禁止发送网络请求。
```

Online Serving

```
1 vllm serve Qwen/Qwen2.5-1.5B-Instruct
```

OpenAI Completions API with vLLM

curl

```
ssh -p 46865 root@region-42.seetacloud.com
X 46865 (ssh)
(APIServer pid=2410) INFO 06-07 12:25:56 [api_server.py:596] Starting vLLM server on http://0.0.0.0:8000
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:37] Available routes are:
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /openapi.json, Methods: HEAD, GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /docs, Methods: HEAD, GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /docs/oauth2-redirect, Methods: HEAD, GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /redoc, Methods: HEAD, GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /tokenize, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /detokenize, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /load, Methods: GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /version, Methods: GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /health, Methods: GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /metrics, Methods: GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/models, Methods: GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /ping, Methods: GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /ping, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /invocations, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/chat/completions, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/chat/completions/batch, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/responses, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/responses/{response_id}, Methods: GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/responses/{response_id}/cancel, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/completions, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/messages, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/messages/count_tokens, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /inference/v1/generate, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /scale_elastic_ep, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /is_scaling_elastic_ep, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /generative_scoring, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/chat/completions/render, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/completions/render, Methods: POST
(APIServer pid=2410) INFO: Started server process [2410]
(APIServer pid=2410) INFO: Waiting for application startup.
(APIServer pid=2410) INFO: Application startup complete.
(EngineCore pid=2488) WARNING 06-07 12:27:45 [jit_monitor.py:103] Triton kernel JIT compilation during inference: _compute_slot_mapping_kerne
l. This causes a latency spike; consider extending warmup to cover this shape/config.
(APIServer pid=2410) INFO: 127.0.0.1:32854 - "POST /v1/completions HTTP/1.1" 200 OK
(APIServer pid=2410) INFO 06-07 12:27:47 [loggers.py:271] Engine 000: Avg prompt throughput: 0.4 tokens/s, Avg generation throughput: 0.7 tok
ens/s, Running: 0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
(APIServer pid=2410) INFO 06-07 12:27:57 [loggers.py:271] Engine 000: Avg prompt throughput: 0.0 tokens/s, Avg generation throughput: 0.0 tok
ens/s, Running: 0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
[]

X 46865 (ssh)
GPU : NVIDIA GeForce RTX 2080 Ti, 1
存储:
系统盘/ : 50% 15G/30G
数据盘/root/autodl-tmp: 53% 27G/50G
+-----+
*注意:
1.系统盘较小请将大的数据存放于数据盘或文件存储中, 重置系统时数据盘和文件存储中的数据不受影响
2.清理系统盘请参考: https://www.autodl.com/docs/qa1/
3.终端中长期执行命令请使用screen等工具开后台运行, 确保程序不受SSH连接中断影响: https://www.autodl.com/docs/daemon/
root@autodl-container-352c469ce5-077f263d:~# curl http://localhost:8000/v1/completions \
-H "Content-Type: application/json" \
-d '{
  "model": "Qwen/Qwen2.5-1.5B-Instruct",
  "prompt": "San Francisco is a",
  "max_tokens": 7,
  "temperature": 0
}'
{"id":"cmpl-b0a120906043ccc2","object":"text_completion","created":1780806465,"model":"Qwen/Qwen2.5-1.5B-Instruct","choices":[{"index":0,"tex
t":" city in the state of California","logprobs":null,"finish_reason":"length","stop_reason":null,"token_ids":null,"prompt_logprobs":null,"p
rompt_token_ids":null,"routed_experts":null},"service_tier":null,"system_fingerprint":"vllm-0.22.1-f3fd6526","usage":{"prompt_tokens":4,"tot
root@autodl-container-352c469ce5-077f263d:~#
```

Python package

```
1 from openai import OpenAI
2
3 # Modify OpenAI's API key and API base to use vLLM's API server.
4 openai_api_key = "EMPTY"
5 openai_api_base = "http://localhost:8000/v1"
6 client = OpenAI(
7     api_key=openai_api_key,
8     base_url=openai_api_base,
9 )
10 completion = client.completions.create(
11     model="Qwen/Qwen2.5-1.5B-Instruct",
12     prompt="San Francisco is a",
13 )
14 print("Completion result:", completion)
```

```
ssh -p 46865 root@region-42.seetacloud.com

46865 (ssh)
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /redoc, Methods: HEAD, GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /tokenize, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /detokenize, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /load, Methods: GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /version, Methods: GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /health, Methods: GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /metrics, Methods: GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/models, Methods: GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /ping, Methods: GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /ping, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /invocations, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/chat/completions, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/chat/completions/batch, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/responses, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/responses/{response_id}, Methods: GET
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/responses/{response_id}/cancel, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/completions, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/messages, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/messages/count_tokens, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /inference/v1/generate, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /scale_elastic_ep, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /is_scaling_elastic_ep, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /generative_scoring, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/chat/completions/render, Methods: POST
(APIServer pid=2410) INFO 06-07 12:25:56 [launcher.py:46] Route: /v1/completions/render, Methods: POST
(APIServer pid=2410) INFO: Started server process [2410]
(APIServer pid=2410) INFO: Waiting for application startup.
(APIServer pid=2410) INFO: Application startup complete.
(EngineCore pid=2488) WARNING 06-07 12:27:45 [jit_monitor.py:103] Triton kernel JIT compilation during inference: _compute_slot_mapping_kerne
l. This causes a latency spike; consider extending warmup to cover this shape/config.
(APIServer pid=2410) INFO: 127.0.0.1:32854 - "POST /v1/completions HTTP/1.1" 200 OK
(APIServer pid=2410) INFO 06-07 12:27:47 [loggers.py:271] Engine 000: Avg prompt throughput: 0.4 tokens/s, Avg generation throughput: 0.7 tok
ens/s, Running: 0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
(APIServer pid=2410) INFO 06-07 12:27:57 [loggers.py:271] Engine 000: Avg prompt throughput: 0.0 tokens/s, Avg generation throughput: 0.0 tok
ens/s, Running: 0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
(APIServer pid=2410) INFO: 127.0.0.1:56884 - "POST /v1/completions HTTP/1.1" 200 OK
(APIServer pid=2410) INFO 06-07 12:29:47 [loggers.py:271] Engine 000: Avg prompt throughput: 0.4 tokens/s, Avg generation throughput: 1.6 tok
ens/s, Running: 0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
(APIServer pid=2410) INFO 06-07 12:29:57 [loggers.py:271] Engine 000: Avg prompt throughput: 0.0 tokens/s, Avg generation throughput: 0.0 tok
ens/s, Running: 0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%

Collecting annotated-types==0.6.0 (from pydantic<3,>=1.9.0->openai)
  Downloading http://mirrors.aliyun.com/pypi/packages/78/b6/6307fbef88d9b5ee7421e68d78a9f162e0da4900bc5f5793f6d3d0e34fb8/annotated_types-0.7.
0-py3-none-any.whl (13 kB)
Collecting pydantic-core==2.46.4 (from pydantic<3,>=1.9.0->openai)
  Downloading http://mirrors.aliyun.com/pypi/packages/5f/97/2aab507d3d00ca626e8e57c1eac6a79e4e5fbcc63eb99733ff55d1717f65/pydantic_core-2.46.4
-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (2.1 MB)
    2.1/2.1 MB 3.3 MB/s eta 0:00:00
Collecting typing-inspection==0.4.2 (from pydantic<3,>=1.9.0->openai)
  Downloading http://mirrors.aliyun.com/pypi/packages/dc/9b/47798a6c91d8bdb567fe2698fe81e0c6b7cb7ef4d13da4114b41d239f65d/typing_inspection-0.
4.2-py3-none-any.whl (14 kB)
Installing collected packages: typing-inspection, pydantic-core, jiter, annotated-types, pydantic, openai
Successfully installed annotated-types-0.7.0 jiter-0.15.0 openai-2.41.0 pydantic-2.13.4 pydantic-core-2.46.4 typing-inspection-0.4.2
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is rec
ommended to use a virtual environment instead: https://pip.pypa.io/warnings/venv
root@autodl-container-352c469ce5-077f263d:~# python3 demo1.py
Completion result: Completion(id='cmpl-a113948aae0de01e', choices=[CompletionChoice(finish_reason='length', index=0, logprobs=None, text=' ci
ty that has seen its fair share of natural disasters. The city was hit hard', stop_reason=None, token_ids=None, prompt_logprobs=None, prompt_
token_ids=None, routed_experts=None)], created=1780806585, model='Qwen/Qwen2.5-1.5B-Instruct', object='text_completion', system_fingerprint='
vllm-0.22.1-f3fd6526', usage=CompletionUsage(completion_tokens=16, prompt_tokens=4, total_tokens=20, completion_tokens_details=None, prompt_t
okens_details=None), service_tier=None, kv_transfer_params=None)
root@autodl-container-352c469ce5-077f263d:~#
```

OpenAI Chat Completions API with vLLM

curl

```
ssh -p 33896 root@connect.bjb1.seetacloud.com

33896 (ssh)
(APIServer pid=3498) INFO: Waiting for application startup.
(APIServer pid=3498) INFO: Application startup complete.
(EngineCore pid=3576) WARNING 06-07 20:16:35 [jit_monitor.py:103] Triton kernel JIT compilation during inference: _compute_slot_mapping_kernel. This causes a latency spike; consider extending warmup to cover this shape/config.
(APIServer pid=3498) INFO: 127.0.0.1:54664 - "POST /v1/completions HTTP/1.1" 200 OK
(APIServer pid=3498) INFO 06-07 20:16:39 [loggers.py:271] Engine 000: Avg prompt throughput: 0.3 tokens/s, Avg generation throughput: 0.4 tokens/s, Running: 0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
(APIServer pid=3498) INFO 06-07 20:16:49 [loggers.py:271] Engine 000: Avg prompt throughput: 0.0 tokens/s, Avg generation throughput: 0.0 tokens/s, Running: 0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
(APIServer pid=3498) INFO: 127.0.0.1:49628 - "POST /v1/completions HTTP/1.1" 200 OK
(APIServer pid=3498) INFO 06-07 20:17:39 [loggers.py:271] Engine 000: Avg prompt throughput: 0.4 tokens/s, Avg generation throughput: 1.6 tokens/s, Running: 0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
(APIServer pid=3498) INFO 06-07 20:17:49 [loggers.py:271] Engine 000: Avg prompt throughput: 0.0 tokens/s, Avg generation throughput: 0.0 tokens/s, Running: 0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
(APIServer pid=3498) INFO: 127.0.0.1:49654 - "POST /v1/chat/completions HTTP/1.1" 200 OK
(APIServer pid=3498) INFO 06-07 20:19:39 [loggers.py:271] Engine 000: Avg prompt throughput: 3.1 tokens/s, Avg generation throughput: 2.4 tokens/s, Running: 0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%

33896 (ssh)
(root) root@autodl-container-70wu030wy8-5a3d0b28:~#
(root) root@autodl-container-70wu030wy8-5a3d0b28:~#
(root) root@autodl-container-70wu030wy8-5a3d0b28:~#
(root) root@autodl-container-70wu030wy8-5a3d0b28:~# curl http://localhost:8000/v1/chat/completions \
-H "Content-Type: application/json" \
-d '{
  "model": "Qwen/Qwen2.5-1.5B-Instruct",
  "messages": [
    {"role": "system", "content": "You are a helpful assistant."},
    {"role": "user", "content": "Who won the world series in 2020?"}
  ]
}'
{"id": "chatcmpl-ae324eec3de2ef0", "object": "chat.completion", "created": 1780834770, "model": "Qwen/Qwen2.5-1.5B-Instruct", "choices": [{"index": 0, "message": {"role": "assistant", "content": "The New York Yankees won the World Series in 2020, defeating the Houston Astros in seven games.", "refusal": null, "annotations": null, "audio": null, "function_call": null, "tool_calls": [], "reasoning": null, "logprobs": null, "finish_reason": "stop", "stop_reason": null, "token_ids": null, "routed_experts": null}}, {"service_tier": null, "system_fingerprint": "vllm-0.22.1-a654ced1", "usage": {"prompt_tokens": 31, "total_tokens": 55, "completion_tokens": 24, "prompt_tokens_details": null, "prompt_logprobs": null, "prompt_token_ids": null, "prompt_text": null, "kv_transfer_params": null}}](root) root@autodl-container-70wu030wy8-5a3d0b28:~#
```

Python package

```
Plain Text |
1 from openai import OpenAI
2 # Set OpenAI's API key and API base to use vLLM's API server.
3 openai_api_key = "EMPTY"
4 openai_api_base = "http://localhost:8000/v1"
5
6 client = OpenAI(
7     api_key=openai_api_key,
8     base_url=openai_api_base,
9 )
10
11 chat_response = client.chat.completions.create(
12     model="Qwen/Qwen2.5-1.5B-Instruct",
13     messages=[
14         {"role": "system", "content": "You are a helpful assistant."},
15         {"role": "user", "content": "Tell me a joke."},
16     ],
17 )
18 print("Chat response:", chat_response)
```



```
ssh -p 33896 root@connect.bjb1.seetacloud.com

33896 (ssh)
(APIServer pid=3498) INFO 06-07 20:16:39 [loggers.py:271] Engine 000: Avg prompt throughput: 0.3 tokens/s, Avg generation throughput: 0.4 tokens/s, Running:
0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
(APIServer pid=3498) INFO 06-07 20:16:49 [loggers.py:271] Engine 000: Avg prompt throughput: 0.0 tokens/s, Avg generation throughput: 0.0 tokens/s, Running:
0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
(APIServer pid=3498) INFO: 127.0.0.1:49628 - "POST /v1/completions HTTP/1.1" 200 OK
(APIServer pid=3498) INFO 06-07 20:17:39 [loggers.py:271] Engine 000: Avg prompt throughput: 0.4 tokens/s, Avg generation throughput: 1.6 tokens/s, Running:
0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
(APIServer pid=3498) INFO 06-07 20:17:49 [loggers.py:271] Engine 000: Avg prompt throughput: 0.0 tokens/s, Avg generation throughput: 0.0 tokens/s, Running:
0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
(APIServer pid=3498) INFO: 127.0.0.1:49654 - "POST /v1/chat/completions HTTP/1.1" 200 OK
(APIServer pid=3498) INFO 06-07 20:19:39 [loggers.py:271] Engine 000: Avg prompt throughput: 3.1 tokens/s, Avg generation throughput: 2.4 tokens/s, Running:
0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
(APIServer pid=3498) INFO 06-07 20:19:49 [loggers.py:271] Engine 000: Avg prompt throughput: 0.0 tokens/s, Avg generation throughput: 0.0 tokens/s, Running:
0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%
(APIServer pid=3498) INFO: 127.0.0.1:54694 - "POST /v1/chat/completions HTTP/1.1" 200 OK
(APIServer pid=3498) INFO 06-07 20:21:09 [loggers.py:271] Engine 000: Avg prompt throughput: 2.4 tokens/s, Avg generation throughput: 1.5 tokens/s, Running:
0 reqs, Waiting: 0 reqs, GPU KV cache usage: 0.0%, Prefix cache hit rate: 0.0%

33896 (ssh)
}
{"id":"chatcmpl-ae324eec3de2e2f0","object":"chat.completion","created":1780834770,"model":"Qwen/Qwen2.5-1.5B-Instruct","choices":[{"index":0,"message":{"role":"assistant","content":"The New York Yankees won the World Series in 2020, defeating the Houston Astros in seven games.", "refusal":null,"annotations":null,"audio":null,"function_call":null,"tool_calls":[],"reasoning":null,"logprobs":null,"finish_reason":"stop","stop_reason":null,"token_ids":null,"routed_experts":null},"service_tier":null,"system_fingerprint":"vllm-0.22.1-a654ced1","usage":{"prompt_tokens":31,"total_tokens":55,"completion_tokens":24,"prompt_tokens_details":null,"prompt_logprobs":null,"prompt_token_ids":null,"prompt_text":null,"kv_transfer_params":null}}](root) root@autodl-container-70wu030wy8-5a3d0b28:~#
#
(root) root@autodl-container-70wu030wy8-5a3d0b28:~#
(root) root@autodl-container-70wu030wy8-5a3d0b28:~#
(root) root@autodl-container-70wu030wy8-5a3d0b28:~# vim chat.py
(root) root@autodl-container-70wu030wy8-5a3d0b28:~# python chat.py
Chat response: ChatCompletion(id='chatcmpl-97d3f956a3913e4b', choices=[Choice(finish_reason='stop', index=0, logprobs=None, message=ChatCompletionMessage(content='Why was the math book sad?\nBecause it had too many problems.', refusal=None, role='assistant', annotations=None, audio=None, function_call=None, tool_calls=[], reasoning=None), stop_reason=None, token_ids=None, routed_experts=None)], created=1780834862, model='Qwen/Qwen2.5-1.5B-Instruct', object='chat.completion', moderation=None, service_tier=None, system_fingerprint='vllm-0.22.1-a654ced1', usage=CompletionUsage(completion_tokens=15, prompt_tokens=24, total_tokens=39, completion_tokens_details=None, prompt_tokens_details=None), prompt_logprobs=None, prompt_token_ids=None, prompt_text=None, kv_transfer_params=None)
(root) root@autodl-container-70wu030wy8-5a3d0b28:~#
```